


Title: Lab 4: Caching Architecture & ETL Flow



- **Node.js API fetches all sensor readings**

The Node.js server acts as the main gateway. Every request for a sensor value flows through it, and it decides whether to serve from Redis or fall back to MongoDB.

- **Cache-Aside & Read-Through both check Redis first**

Before touching MongoDB, both strategies try to pull the value from Redis to save time. Cache-Aside loads into Redis only after a miss; Read-Through automatically fills Redis whenever data is missing.

- **MongoDB remains the system of record**

Even though Redis speeds everything up, MongoDB still stores the actual data long-term. Redis only holds the “fast copies.”

- **TTL ensures stale values get refreshed automatically**

TTL-based caching forces entries to expire after a few seconds. Once expired, the next request triggers a clean reload from MongoDB.

- **ETL pipeline stores raw data in MongoDB, Redis caches only on demand**

Incoming sensor readings are written directly to MongoDB. Redis is not preloaded—values only appear in Redis when your app explicitly reads them and caches them afterward.

— Lessons Learned

- **Built and tested multiple caching strategies in one pipeline**

Implementing Cache-Aside, Read-Through, and TTL helped me understand how each pattern behaves differently and where each one is useful.

- **Redis massively reduces response time**

Before caching, MongoDB queries averaged around 55 ms. After caching, Redis hits dropped to around 1 ms. The performance difference was obvious and measurable.

- **Learned how to structure clean Node.js cache helper functions**

Breaking the logic into small, reusable functions (like `readThrough()` and `cache setters`) made the server easier to debug and extend.

- **Understood how TTL affects real-time performance**

Seeing requests slow down immediately after a TTL expiry (fresh reload) and then speed up again once Redis is populated helped reinforce how expiration-based caching works.

- **Saw how important cache-DB consistency is**

If the database changes but Redis doesn't get updated, you serve stale data. Handling this consistently is a real design challenge.

- **Learned core distributed system ideas: freshness, invalidation, and cache coherence**

This project made these concepts real instead of theoretical—especially when I had to clear or reset Redis to simulate fresh reads.

What Worked Well

- **Redis integration worked smoothly once the setup was stable**

After correctly configuring the Redis client, the connection stayed solid and data retrieval was instant.

- **Cache-Aside gave consistent performance boosts**

The HIT vs MISS behavior was very clear. Once cached, the API felt significantly faster.

- **TTL caching behaved exactly as expected**

Entries expired right on schedule, and after expiry the system correctly reloaded and recached the data.

- **MongoDB queries were stable and predictable**

The database always returned clean, structured documents, making the pipeline reliable.

Challenges

- **Route conflicts originally caused “Cannot GET /cache-aside/...” errors**

This happened because multiple versions of the same route were defined in the server file, causing Express to break. Fixing it required cleaning duplicates.

- **Multiple Node.js servers were running, causing port 3000 to be locked**

Before restarting, I had to manually kill the process using lsof and kill -9, otherwise nothing worked.

- **Redis had to be flushed repeatedly to test proper cache behavior**

To simulate fresh misses, I had to manually clear Redis keys every time using FLUSHALL or DEL. Forgetting this produced confusing results.

- **TTL testing required precise timing**

If I didn't wait the correct number of seconds, I would get a Redis HIT instead of an expiry reload, which made testing slightly annoying.