

# Architecture & ETL Data Flow (Redis → MongoDB)

- DATA SOURCE INGESTED  
USING ETL SCRIPTS  
(EXTRACT, TRANSFORM,  
LOAD)

- REDIS ACTS AS A HIGH-SPEED CACHING LAYER FOR RAPID DATA RETRIEVAL

- MONGODB DATA LAKE STORES PROCESSED AND HISTORICAL DATA

- DASHBOARD TOOLS CONNECT FOR ANALYTICS AND VISUALIZATION

# Lessons Learned from Extending Cached Database Pipeline



- Redis significantly improved query response times and scalability.



- MongoDB's flexible schema allowed integration of diverse data sources.



- ETL automation ensured consistent, reliable data ingestion.



- Monitoring Redis TTL and eviction policies was critical for accuracy.



- Streamlined analytics improved data-driven decision-making.

- What Worked Well:
  - Caching reduced latency and optimized data retrieval.
  - ETL automation improved consistency and reduced manual errors.
- What Did Not Work Well:
  - Cache invalidation required manual intervention at times.
  - Synchronizing Redis and MongoDB was complex during heavy loads.
  - Memory management in Redis needed careful tuning.
- Future Improvements:
  - Introduce automated cache refresh logic.
  - Implement monitoring dashboards for better system insights.

Challenges  
What  
Worked  
Well?  
What Did  
Not Work  
Well?