

# Redis–MongoDB Data Pipeline Architecture

Data Sources → ETL → Redis Cache → MongoDB Atlas → Reporting/Dashboard

Flow:

- Data collected from IoT/API sensors
- ETL transforms data to JSON format
- Redis stores frequently accessed readings
- MongoDB stores long-term sensor history
- Cache refresh & TTL maintain data freshness

# Lessons Learned

- Redis caching significantly reduced latency ( $\sim 100 \text{ ms} \rightarrow \sim 5 \text{ ms}$ )
- TTL & Cache-Aside strategies balance performance and freshness
- Integration of Redis + MongoDB improved scalability
- Logging & benchmarking helped query optimization
- Cache invalidation policies prevent stale data

# Challenges & Outcomes

What worked well:

- Stable Redis connection
- Fast reads under high load
- Streamlined ETL flow into MongoDB

What didn't work well:

- Cache eviction & TTL tuning needed trial & error
- Occasional MongoDB delays on cache reload
- Concurrent updates required additional logic

Outcome:

Reliable hybrid caching + database pipeline supporting real-time analytics and historical storage