

Architecture & ETL Data Flow (Redis → MongoDB)

- **Architecture Overview:**
 - - Sensors / App data → Node.js API → Redis cache → MongoDB Atlas (AgriDB.readings)
 - - Redis holds hot data for fast retrieval (Cache-Aside, Read-Through, TTL)
 - - MongoDB stores persistent data for long-term analytics
- **ETL Flow:**
 - 1. Extract – Node.js collects readings via API
 - 2. Transform – Validate JSON, add metadata (_ingestedAt, _source, meta.author)
 - 3. Load – Insert into MongoDB; update Redis cache
- Key Message: Redis = Speed Layer, MongoDB = Storage Layer

Lessons Learned from Extending the Cached Database Pipeline

- Integrating Redis reduced latency from ~120 ms → ~20 ms (after caching)
- Learned Cache-Aside, Read-Through, and TTL expiration patterns
- Improved ETL pipeline with metadata tracking and validation
- Thunder Client testing showed clear cache-hit and TTL behavior

- Key Takeaway: Combining caching logic with ETL ensures speed and data freshness.

Challenges, What Worked Well, and What Did Not Work Well

- **Worked Well:**
 - - Reliable Redis + MongoDB connections
 - - Cache-Aside pattern gave visible performance gain
 - - TTL handled auto-expiration effectively
- **Challenges:**
 - - Manual cache invalidation after DB updates
 - - Debugging Redis connection errors
 - - Measuring TTL expiry precisely
- **Next Steps:**
 - Automate invalidation using Pub/Sub
 - Implement Write-Through caching
 - Add visualization for cache performance metrics